

Programmed By : Rithik Tripathi

[Connect with me on LinkedIn \(https://www.linkedin.com/in/rithik-tripathi-data-scientist/\)](https://www.linkedin.com/in/rithik-tripathi-data-scientist/)

In this notebook, We will implement Logistic Regression and Understand how we can CHANGE THRESHOLD for classification

## Logistic Regression Demonstration

In [4]:

```
1 import pandas as pd
2 import numpy as np
```

In [70]:

```
1 df = pd.read_csv("Dataset/KNN/data_cleaned.csv")
2 df.head()
3
4 # using Titanic Survival Dataset which has been already cleaned and encoded for usage
```

Out[70]:

	Survived	Age	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	SibSp_0	SibSp_1	...	Parch_0	Parch_1	Parch_2	Parch_3	Parch_4	P
0	0	22.0	7.2500	0	0	1	0	1	0	1	...	1	0	0	0	0	
1	1	38.0	71.2833	1	0	0	1	0	0	1	...	1	0	0	0	0	
2	1	26.0	7.9250	0	0	1	1	0	1	0	...	1	0	0	0	0	
3	1	35.0	53.1000	1	0	0	1	0	0	1	...	1	0	0	0	0	
4	0	35.0	8.0500	0	0	1	0	1	1	0	...	1	0	0	0	0	

5 rows × 25 columns

## Separating Independent & Target Features

In [6]:

```
1 x = df.drop(['Survived'], axis=1)
2 y = df['Survived']
3 x.shape, y.shape
```

Out[6]:

((891, 24), (891,))

In [10]:

```
1 from sklearn.model_selection import train_test_split
2 x_train, x_test, y_train, y_test = train_test_split(x,y, random_state=9, stratify=y)
3
4 x_train.shape, x_test.shape
```

Out[10]:

((668, 24), (223, 24))

## Scaling Features

using Min Max scalar

In [16]:

```
1 from sklearn.preprocessing import MinMaxScaler
2 scaler = MinMaxScaler()
3
4 train_x_scaled = scaler.fit_transform(x_train)
5 train_x_scaled = pd.DataFrame(train_x_scaled, columns=x_train.columns)
6 train_x_scaled.head()
```

Out[16]:

	Age	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	SibSp_0	SibSp_1	SibSp_2	...	Parch_0	Parch_1	Parch_2	Parch_3	Parch_4
0	0.044986	0.032596	0.0	0.0	1.0	1.0	0.0	0.0	1.0	0.0	...	0.0	1.0	0.0	0.0	0
1	0.346569	0.025374	0.0	1.0	0.0	1.0	0.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0
2	0.367921	0.030254	0.0	0.0	1.0	1.0	0.0	0.0	1.0	0.0	...	1.0	0.0	0.0	0.0	0
3	0.258608	0.015127	0.0	0.0	1.0	1.0	0.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0
4	0.233476	0.019852	0.0	0.0	1.0	0.0	1.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0

5 rows × 24 columns

In [17]:

```
1 test_x_scaled = scaler.transform(x_test)
2 test_x_scaled = pd.DataFrame(test_x_scaled, columns=x_test.columns)
3 test_x_scaled.head()
```

Out[17]:

	Age	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	SibSp_0	SibSp_1	SibSp_2	...	Parch_0	Parch_1	Parch_2	Parch_3	Parch_4
0	0.648153	0.059532	1.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0
1	0.296306	0.096626	1.0	0.0	0.0	1.0	0.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0
2	0.346569	0.051822	1.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0
3	0.585323	0.049943	1.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0
4	0.258608	0.016461	0.0	0.0	1.0	0.0	1.0	1.0	0.0	0.0	...	1.0	0.0	0.0	0.0	0

5 rows x 24 columns

Implementing Logistic Regression

In [1]:

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.metrics import f1_score
```

In [26]:

```
1 # creating instance of Logistic Regression class
2 LogisticReg = LogisticRegression(max_iter=500)
3
4 # Training / Fitting the model
5 LogisticReg.fit(x_train, y_train)
```

Out[26]:

LogisticRegression(max\_iter=500)

We are aware of the fact that each feature has a coefficient associated with its equation.

Lets check them out

In [28]:

```
1 pd.DataFrame({
2     'Feature': x_train.columns,
3     'Coefficients': LogisticReg.coef_.flatten()
4 }).sort_values(by = 'Coefficients')
```

Out[28]:

	Feature	Coefficients
6	Sex_male	-1.291244
4	Pclass_3	-0.938834
10	SibSp_3	-0.899003
18	Parch_4	-0.801684
19	Parch_5	-0.698865
13	SibSp_8	-0.452825
12	SibSp_5	-0.232781
11	SibSp_4	-0.193704
23	Embarked_S	-0.159248
0	Age	-0.031875
22	Embarked_Q	-0.002438

We could clearly see that Probability of Surviving is Minimal with sex\_male, 3rd class passengers and so on... (having smallest coefficient) and Maximum with sex\_female, 1st class passengers and so on ... (having highest coefficient)

## Predictions

In [37]:

```
1 train_pred = LogisticReg.predict(x_train)
2 train_pred[:10]
3
4 # these are the hard tags for survival (sliced out only 10 predictions for better visibility)
```

Out[37]:

```
array([1, 1, 1, 1, 0, 0, 1, 0, 0, 0], dtype=int64)
```

### Predicting over Train set

In [39]:

```
1 # calculating f1_score for train set
2 train_f1_score = f1_score(y_train, train_pred)
3 print('Training F1 Score : ',train_f1_score)
4
5 # remeber this score is with the default 50% Threshold
```

Training F1 Score : 0.7398373983739838

### Predicting over Test set

In [40]:

```
1 test_pred = LogisticReg.predict(x_test)
2 test_pred[:10]
```

Out[40]:

```
array([0, 1, 0, 0, 0, 0, 0, 1, 1, 1], dtype=int64)
```

In [42]:

```
1 # calculating f1_score for test set
2 test_f1_score = f1_score(test_pred, y_test)
3 print('Test F1 Score : ',test_f1_score)
4
5 # remeber this score is with the default 50% Threshold
```

Test F1 Score : 0.7349397590361445

**We could say model is not overfitting here and the scores are descent enough**

As seen earlier in this notebook , the predictions predicted by .predict() are Hard Tags

In order to generate the Soft Tags , we will use .predict\_proba()

## Predicting Soft Tags

In [53]:

```
1 train_pred_soft = LogisticReg.predict_proba(x_train)
2 train_pred_soft[:10]
3
4 # the output here is the COLUMNS equal to the number of CLASSES in TARGET
5 # and EACH COLUMN CONTAINS PROBABILITY OF THAT DATA POINT BELONGING TO THAT CLASS
6
7 # the below data shows the probability of each data point for 1st 10 records belonging to class 0 & 1 respectively
```

Out[53]:

```
array([[0.16079452, 0.83920548],
       [0.21442751, 0.78557249],
       [0.35700362, 0.64299638],
       [0.33732195, 0.66267805],
       [0.8810176 , 0.1189824 ],
       [0.88812343, 0.11187657],
       [0.22680276, 0.77319724],
       [0.91303642, 0.08696358],
       [0.60869506, 0.39130494],
       [0.88404956, 0.11595044]])
```

In [51]:

```
1 # we can check and validate the sum across the axis will always be 1 (total probability)
2 train_pred_soft.sum(axis=1)[:10]
```

Out[51]:

```
array([1., 1., 1., 1., 1., 1., 1., 1., 1., 1.])
```

The best part about having probabilities is that we can set our own Thresholds to segregate between classes ¶

In [54]:

```
1 # seggeregating survived column from probabilities
2 train_pred_soft_survived = train_pred_soft[:, 1]
3 train_pred_soft_survived[:10]
```

Out[54]:

```
array([0.83920548, 0.78557249, 0.64299638, 0.66267805, 0.1189824 ,
       0.11187657, 0.77319724, 0.08696358, 0.39130494, 0.11595044])
```

In [55]:

```
1 # setting up threshold of 0.55 and separating
2
3 for i in range(len(train_pred_soft_survived)):
4
5     if(train_pred_soft_survived[i] > 0.55):
6         train_pred_soft_survived[i] = 1
7     else:
8         train_pred_soft_survived[i] = 0
```

In [57]:

```
1 # checking how they Look now
2 train_pred_soft_survived[:10]
```

Out[57]:

```
array([1., 1., 1., 1., 0., 0., 1., 0., 0., 0.])
```

In [59]:

```
1 # calculating F1 Score
2
3 train_f1_score_soft = f1_score(y_train, train_pred_soft_survived)
4 print('Training F1 Score (soft) : ',train_f1_score_soft)
5
6 # remeber this score is with the custom 55% Threshold
7
8 # notice the change in accuracy
```

```
Training F1 Score (soft) : 0.7375
```

## Creating Confusion Matrix

In [61]:

```
1 from sklearn.metrics import confusion_matrix, log_loss
2
3 confusion_mat = confusion_matrix(y_test, test_pred)
4 print(confusion_mat)
```

```
[[118 19]
 [ 25 61]]
```

## Generating Classification Report

In [62]:

```
1 from sklearn.metrics import classification_report as report
2 print(report(y_test, test_pred))
```

	precision	recall	f1-score	support
0	0.83	0.86	0.84	137
1	0.76	0.71	0.73	86
accuracy			0.80	223
macro avg	0.79	0.79	0.79	223
weighted avg	0.80	0.80	0.80	223

```
1 Support => shows that how many data points are belonging to that particular class
2 macro avg => shows the average of above two entities irrespective of class
3 weighted avg => average calculated by their respective weights (weights are the points available in the support of that class)
```

Visualizing the coefficients

In [63]:

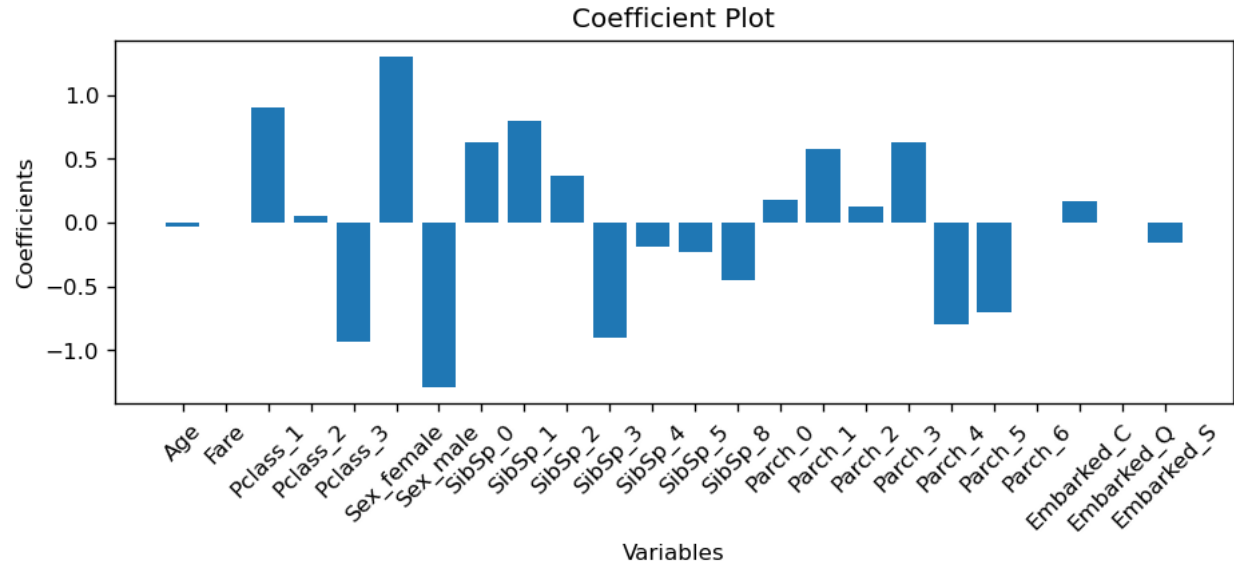
```
1 LogisticReg.coef_
```

Out[63]:

```
array([[ -0.03187542,  0.00261003,  0.89636692,  0.05059092, -0.93883372,
         1.29936852, -1.2912444 ,  0.62758533,  0.79778733,  0.36106487,
        -0.89900305, -0.19370445, -0.23278051, -0.45282541,  0.17803249,
         0.57647513,  0.12804463,  0.62612123, -0.80168393, -0.69886542,
         0.          ,  0.16981082, -0.00243845, -0.15924825]])
```

In [67]:

```
1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize = (9,3), dpi=120, facecolor='w', edgecolor='b')
4 x = x_train.columns
5 y = LogisticReg.coef_.reshape(-1)
6 plt.bar(x,y)
7 plt.xlabel('Variables')
8 plt.ylabel('Coefficients')
9 plt.title('Coefficient Plot')
10 plt.xticks(rotation = 45)
11 plt.show()
```



In [ ]:

```
1
```